Turn every lost sale into a coached next action

Offer Blocker Intelligence on Databricks

Business outcome: convert unstructured sales calls into a governed queue of price, fee, contract, availability, and promotion barriers—fast enough for a manager to intervene while the opportunity is still actionable.

Decision requested: sponsor a measured pilot with Sales Operations and Inside Sales.

Superior Plus Propane · Evaluation submission · October 2026

The business loses signal between the call and the coaching moment

Sales calls contain the reason an offer stalled, but the signal is trapped in long transcripts.

Today	Business consequence
Managers sample calls manually	Coverage is partial and selection is inconsistent
Blockers are coded after the fact	Follow-up arrives after buyer intent cools
Evidence lives in prose	Pricing, offer, and enablement teams cannot aggregate patterns
Coaching actions live outside analytics	Ownership and outcomes are hard to audit

The opportunity is not “better transcription.” It is a closed loop from **conversation → blocker → owner → action → outcome**.

One governed flow turns calls into action

1. **Ingest** transcripts incrementally with Auto Loader and deduplicate by call ID.
2. **Contextualize** each call against a synthetic CRM opportunity and ordered dialogue.
3. **Classify** six offer-blocker categories with Databricks AI Functions.
4. **Explain** every finding with disposition, confidence, and quoted evidence.
5. **Explore** patterns in an AI/BI dashboard and curated Genie space.
6. **Act** in the coaching app: assign a case, capture notes, agree actions, and retain an audit trail in Lakebase.

Governance by design: Unity Catalog tables, evidence-bearing output, PII-masked summaries, service-principal write-back, and no customer data in this submission.

The reference run proves the pipeline emits decision-ready output

The executed, pseudonymized reference output contains:

Measure	Observed result
Total output rows	579
Classified 4A–4F findings	273
Hard blockers	47
Friction findings	42
Resolved issues	27
Taxonomy coverage	All 6 codes

Example: a customer rejects a standard installation fee and leaves to shop elsewhere → **4B Ancillary Fees · hard_blocker · High confidence**.

The build also passed **61/61 application tests**, TypeScript compilation, lint, formatting, and production bundling. Full text evidence is committed under `submission/`.

Modeled pilot economics: reclaim 108 manager hours each month

Illustrative planning model—validate with pilot telemetry.

Assumption	Value
Opportunities reviewed per month	500
Manual review and coding	15 min/opportunity
AI-assisted triage	2 min/opportunity
Time reclaimed	**108 hours/month**
Review effort reduction	**87%**

Formula: `500 × (15 − 2) ÷ 60 = 108 hours/month`.

At a loaded manager cost of **$75/hour**, that is approximately **$8.1K/month** or **$97.5K/year** of capacity before counting recovered deals. These are modeled benefits, not realized claims.

Value for the executive sponsor

Revenue visibility

- Quantify which offer mechanics suppress conversion instead of relying on anecdotes.
- Detect shifts in rate, fee, contract, serviceability, and promotion objections.

Operating leverage

- Move from sampled call review toward systematic coverage.
- Give managers evidence-first cases rather than another dashboard to interpret.

Governance and trust

- Preserve source evidence, generated SQL, manager feedback, and action history.
- Keep analytical read paths and operational write-back under Databricks governance.

Sponsor KPI: **cost per reviewed opportunity**, **coverage rate**, and **conversion lift on coached opportunities**.

Value for the Sales Operations domain owner

Every day

- Filter the queue to real blockers, overdue follow-ups, rep, or blocker code.
- Open a case with the exact evidence and AI confidence in context.
- Assign coaching, record the handling score, and track agreed actions.
- Ask Genie follow-up questions and preserve answer provenance.

Every week

- Review hard-blocker and friction trends by category.
- Separate product/offer gaps from issues the representative resolved.
- Feed validated false positives back into prompt and taxonomy calibration.

Domain-owner KPI: **time to first coaching action**, **action completion**, **finding acceptance rate**, and **repeat-blocker rate**.

A four-week pilot converts modeled value into measured value

Week	Scope	Exit criterion
1	Baseline manual review time, select one team, confirm taxonomy	Agreed KPI definitions and safe data contract
2	Run pipeline and manager queue in shadow mode	≥90% records processed; evidence present on every assigned finding
3	Managers coach from the queue and rate findings	≥70% manager acceptance; median action time measured
4	Compare coached cohort and baseline; tune thresholds	Go/no-go recommendation with measured capacity and conversion signals

Guardrails: synthetic data for demos, least-privilege access, human review before action, and no claim of causality without a controlled comparison.

Decision: fund the pilot and hold it to buyer KPIs

Approve a one-team, four-week pilot owned jointly by Sales Operations and Inside Sales leadership.

Success means:

- **Coverage:** at least 90% of eligible opportunities analyzed.
- **Speed:** median time from call to coaching action below one business day.
- **Trust:** at least 70% of surfaced blockers accepted by managers.
- **Execution:** at least 80% of agreed actions completed by due date.
- **Economics:** measured review effort reduced by at least 60%.

The build is ready for evaluation: deployable code, dashboard and Genie definitions, tested coaching app, text-readable execution evidence, and a data-safe synthetic demo.